

XRA-GS: X-ray Attenuation-Aligned Gaussian Splatting for Sparse Tomographic View Synthesis

Abstract

*3D Gaussian Splatting has been extended to X-ray novel view synthesis, yet existing methods retain the gradient-driven densification rules from natural-image 3DGS without modification. Under sparse CT acquisition, these rules cause Gaussians to accumulate near high-contrast anatomical boundaries while interior tissues that determine X-ray attenuation remain under-represented. We present **XRA-GS**, an X-ray Attenuation-Aligned Gaussian Splatting framework that corrects this structural imbalance by redesigning the Gaussian evolution pipeline. XRA-GS introduces three components: a support-profile seeding strategy that initializes Gaussians from the coarse attenuation support estimated by FDK reconstruction; a geometry-aware pruning step that removes Gaussians concentrated near anatomical boundaries; and an adaptive density modulation step that applies position-dependent density refinement throughout training. Experiments on five organs under a sparse-view CT protocol demonstrate that XRA-GS achieves state-of-the-art novel view synthesis quality against six representative baselines from both visible-light sparse-view synthesis and X-ray CT methods.*

1. Introduction

Computed tomography (CT) is an essential imaging technique for noninvasively examining the internal structure of an object, and is widely used in clinical diagnosis, image-guided intervention, and industrial inspection [FDK84; AK84; WYD08]. During a CT acquisition, an X-ray source rotates around the object while a detector records multi-angle attenuation measurements, from which cross-sectional images or projections at unseen angles can be re-

covered through tomographic reconstruction. Reducing the number of acquired views is highly desirable, since it directly lowers radiation dose, shortens scan time, and broadens applicability to patients who cannot sustain long, stable acquisitions [LMC*23; WHL*24]. In this work, we study *Sparse Tomographic View Synthesis*: from only a handful of X-ray projections, the model must predict the line-integrated attenuation observed at unseen acquisition angles, so that each synthesized projection remains consistent with the Beer–Lambert imaging ge-

ometry rather than merely visually similar to the inputs. A reliable solution supplies auxiliary views for anatomical localization, lesion-range inspection, and intraoperative planning under severely limited sampling, where conventional analytical reconstruction breaks down [WN19].

Visible-light and X-ray imaging are governed by fundamentally different rendering equations, and this difference dictates where representation capacity should reside. In visible-light 3DGS, a pixel is formed by front-to-back α -compositing of an emission-absorption volume rendering integral [KKLD23; Max95]: the accumulated transmittance decays multiplicatively along the ray, so for opaque scenes the integrand is dominated by Gaussians near the first-hit surface while contributions from deeper Gaussians are progressively masked. The surface-clustered Gaussian layout that emerges under α -compositing is therefore a direct consequence of this transmittance-induced concentration, and is consistent with the surface-aligned geometry that surface-oriented 3DGS variants exploit explicitly [GL24; HYC*24]. In X-ray imaging, however, the detector reading follows the Beer–Lambert law [KS01]: after taking the logarithm, the recorded quantity becomes a linear, order-independent line integral of the attenuation coefficient along the ray, with no transmittance decay and no occlusion. Consequently, every Gaussian on the path contributes on equal footing, and capacity must be distributed *along the entire ray path* rather than collapsed onto any single layer. Existing X-ray Gaussian methods such as X-Gaussian, R2-Gaussian, X-Field, and DGR have correctly replaced α -compositing with a radiative line-integral renderer [CLW*24; ZLC*24; WTW*25; WLW*25], and sparse-view 3DGS pipelines such as FSGS, CoR-GS, and DNGaussian further compensate for missing observations through depth or co-

regularization priors [ZFJW24; ZLY*24; LZB*24]. Yet all of them inherit a Gaussian evolution rule that was originally designed for surface-centric α -compositing [KKLD23; RNSZ24].

Crucially, the issue is not gradient-driven densification itself. Splitting or cloning Gaussians where the projection-error gradient is large is a sound idea, and our method retains exactly this mechanism. The issue is that under a line-integral renderer, the residual gradient is still concentrated at high-attenuation-contrast interfaces, such as bone–soft-tissue boundaries, which almost coincide with the high-color-contrast interfaces that drive visible-light densification. A naïve port therefore keeps spawning Gaussians at anatomical boundaries even though the rendering equation now asks for capacity *along* the path. Under sparse views, this bias surfaces not as insufficient coverage but as a structural *capacity misallocation*: a substantial fraction of Gaussians is spent on boundary-driven redundancy, while interior path segments that determine the attenuation integral remain under-represented. The top row of Fig. 1 shows the natural fit between α -compositing and surface-clustered Gaussians; the bottom row shows that when the same evolution rule is reused under line-integral rendering, capacity still collapses to boundaries and leaves the deep path empty. Fixing this requires *augmenting*, not discarding, gradient-driven densification with attenuation-aligned decisions about where to seed, what to prune, and how to refine.

Motivated by this analysis, we propose **XRA-GS**, an X-ray Attenuation-Aligned Gaussian Splatting framework for Sparse Tomographic View Synthesis. Rather than appending auxiliary branches to a conventional densification backbone, **XRA-GS** keeps the gradient-driven backbone of standard 3DGS and adds three attenuation-aligned stages that

act on *where* Gaussians are seeded, retained, and refined, as illustrated in Fig. 1. Support-Profile Seeding (SPS) anchors initialization to a coarse FDK attenuation support [FDK84], so that Gaussians populate physically meaningful X-ray paths instead of starting from a contrast-biased prior; Geometry-aware Pruning (GAP) recycles Gaussians that accumulate near high-contrast anatomical interfaces during training, draining boundary-driven redundancy at its source; and Adaptive Density Modulation (ADM) applies position-dependent online refinement so that retained Gaussians carry interior attenuation rather than uniformly inflated density. Across diverse organs and view budgets, this attenuation-aligned evolution reallocates representation capacity from boundary redundancy to attenuation-relevant interior structures.

Our main contributions are summarized as follows:

- We identify *capacity misallocation* as the failure mode when sparse-view X-ray Gaussian methods inherit a visible-light evolution rule under a line-integral renderer: gradient-driven densification keeps clustering Gaussians at attenuation boundaries while interior path segments that determine the line integral remain starved of capacity.
- We propose **XRA-GS**, an X-ray Attenuation-Aligned Gaussian Splatting framework that augments standard gradient-driven densification with three coordinated attenuation-aligned stages—Support-Profile Seeding (SPS) for path-anchored initialization from coarse FDK support, Geometry-aware Pruning (GAP) for boundary-redundancy recovery, and Adaptive Density Modulation (ADM) for position-dependent online refinement—without modifying the X-ray renderer.

- Extensive experiments on multi-organ sparse-view CT show that **XRA-GS** delivers consistent gains in projection fidelity over state-of-the-art X-ray Gaussian and sparse-view 3DGS base-lines, while reducing the number of Gaussians and training cost; ablations confirm that SPS, GAP, and ADM each contribute complementary capacity reallocation.

2. Related Work

2.1. Gaussian Splatting

3D Gaussian Splatting (3DGS) represents a scene with a set of anisotropic Gaussians and renders novel views in real time through differentiable tile-based rasterization [KKLD23]. Building on this explicit representation, subsequent work has expanded along several directions: surface-aligned variants that recover geometrically accurate meshes from Gaussians [HYC*24; GL24]; dynamic and deformable extensions for time-varying scenes [WRW*24; YGZ*24]; dense visual SLAM systems built on top of the Gaussian representation [YQX*24]; generative pipelines that lift 3DGS into 3D content creation [TRZ*24]; and compactness/efficiency redesigns that reduce the Gaussian count or stabilize rendering across scales [NMKW24; LYX*24; YCH*24]. Across all these directions, the underlying evolution rule remains the gradient-driven clone-and-split densification together with spherical-harmonics view-dependent color, both of which are tailored to dense-view, surface-reflectance scenes. When either assumption is violated—inputs become sparse, or the imaging model switches from surface radiance to line-integral attenuation—this rule starts to misallocate Gaussian capacity, and the next two subsections examine these two regimes in turn.




Figure 1: Visible-light vs. X-ray Gaussian rendering and the capacity gap they impose. (a) Visible-light α -compositing concentrates capacity on visible surfaces, where front Gaussians occlude back ones. (b) X-ray line integrals weight every Gaussian along the ray equally, so naïve gradient-driven densification still clusters at attenuation boundaries; **XRA-GS** realigns the evolution rule through SPS, GAP, and ADM to anchor capacity along the path.

2.2. Sparse View Synthesis

Neural Radiance Fields (NeRF) recover a continuous volumetric field through coordinate MLPs and volume rendering [MST*21], but degenerate sharply when only a handful of input views are available. NeRF-side sparse-view methods inject external priors to constrain the under-determined optimization: depth supervision from sparse structure-from-motion points [DLZR22; WCLL23], semantic consistency from pretrained vision encoders [JTA21], and geometric/appearance regularization on sampled patches together with frequency-domain annealing [NBM*22; YPW23]. Sparse-view 3DGS work inherits the same diagnosis but redirects regularization onto the explicit Gaussian set. Per-scene optimizers diverge in how they explain failure: FSGS attributes degradation to insufficient coverage and augments densification with proximity-

guided unpooling and depth priors [ZFJW24]; CoR-GS treats it as unstable optimization and introduces dual-field co-regularization with co-pruning [ZLY*24]; DNGaussian injects geometric priors through depth normalization [LZB*24]; and DropGaussian counters overfitting through stochastic Gaussian dropout [PRK25]. A parallel feed-forward line moves geometry completion into large-scale pretraining and synthesizes Gaussians from image features in a single forward pass [CLTS24; CXZ*24; XPW*25; LFY*25]. Beyond these regularizers, plane-decomposed spatial encodings such as K-Planes provide continuous positional context at low cost [FYT*23], and we reuse this encoding inside ADM for position-dependent density refinement. Whether the remedy is depth priors, co-regularization, dropout, or feed-forward inference, all of these methods retain an error/gradient-driven

densification backbone that prioritizes high-contrast surfaces. This backbone implicitly assumes visible-light surface radiance; as the next subsection makes explicit, the same logic systematically misallocates capacity once the renderer becomes a line integral over attenuation.

2.3. Tomographic Novel View Synthesis

Sparse-view tomographic novel view synthesis aims to render unseen X-ray projections, and to recover the underlying density volume, from a handful of input projections, thereby reducing radiation exposure relative to the hundreds of views required by classical CT pipelines [WYD08; WN19]. Analytical reconstructions such as FDK invert the Radon transform in closed form but produce severe streak artifacts under sparse views [FDK84], while iterative algebraic solvers such as SART suppress artifacts at the cost of long iterations [AK84]. Learning-based methods address these limits along two lines. The implicit family adapts neural attenuation fields to X-ray imaging: IntraTomo casts tomography as self-supervised sinogram synthesis with a coordinate field [ZIL*21]; NAF accelerates inference through hash-encoded attenuation fields [GWS*24]; SAX-NeRF couples a transformer with structural line-dependency modeling [CWY*24]; and Geometry-Aware Attenuation Learning shifts toward population-level encoder-decoder inference [LFL*25]. The explicit family redesigns Gaussians around X-ray attenuation: X-Gaussian introduces radiative Gaussians together with a differentiable X-ray rasterizer [CLW*24]; R2-Gaussian aligns the representation with reconstruction consistency through integration-bias correction and differentiable voxelization [ZLC*24]; X-Field couples material-aware ellipsoids with explicit segment-length modeling for physics-aligned attenuation [WTW*25]; and a further set of extensions tackles discretized voxel reconstruction [WLG*25],

artifact-suppression via graph-structured gradients [YMS*25], 4D dynamic CT [YCZ*25], and slice-wise volumetric reconstruction [Kua*25]. The latter four target imaging regimes that fall outside the sparse-view novel-view projection protocol our experiments evaluate, and we therefore omit them from the main quantitative comparison. Our work belongs to this explicit radiative-Gaussian line, but its contribution is not another physical branch. Even with a physically aligned renderer, sparse-view CT still suffers if Gaussian capacity allocation inherits the boundary-chasing logic of visible-light densification. This misallocation operates at three distinct stages: *initialization* that ignores the X-ray attenuation support, *boundary over-densification* during training, and *undifferentiated point-wise density updates* in later refinement. SPS, GAP, and ADM each address one of these stages in turn, forming an attenuation-aligned Gaussian evolution pipeline inside **XRA-GS** rather than a modified rendering branch.

3. Method

Figure 1 already summarizes the core distinction at the pipeline level: **XRA-GS** does not keep stacking auxiliary branches on top of an existing densification backbone, but instead reorders how Gaussians are initialized, reclaimed, and refined. We now turn to the implementation level and show how this attenuation-aligned evolution pipeline is instantiated by SPS, GAP, and ADM on top of the radiative Gaussian backbone of R2-Gaussian.

3.1. Problem Formulation

Let the sparse-view observation set be $\mathcal{P} = \{(I_v, \pi_v)\}_{v=1}^V$, where I_v denotes the v -th X-ray projection and π_v denotes its corresponding acquisition geometry. Given a coarse FDK reconstruction volume V_{FDK} obtained from these observations, our goal is to

learn an explicit radiative Gaussian representation in a per-case optimization setting so that unseen-angle target projections can be synthesized from only a few observed projections. Although this task is naturally coupled with volumetric density recovery, our primary optimization target remains novel-view quality in the projection domain.

3.2. 3DGS and X-ray Rendering Background

To clarify that our problem is not simply “adding another prior,” we first contrast the pixel formation mechanisms of visible-light 3DGS and X-ray projection. In standard visible-light 3DGS, the color of a pixel is typically accumulated through front-to-back compositing of multiple splats ordered along the viewing ray:

$$\mathbf{C}(\mathbf{r}) = \sum_{i=1}^{N_r} T_i \alpha_i \mathbf{c}_i, \quad T_i = \prod_{j<i} (1 - \alpha_j), \quad (1)$$

where $\mathbf{r}$ is the camera ray, and α_i and $\mathbf{c}_i$ denote the opacity and radiance attributes of the i -th splat. The core intuition is that a pixel is formed by the accumulation of local surface contributions.

In X-ray imaging, however, a detector pixel corresponds more closely to the remaining energy of a single ray after traversing internal materials. Its physical measurement can be written as

$$I(\mathbf{r}) = I_0 \exp\left(-\int_{\mathbf{r}} \mu(\mathbf{x}) d\mathbf{l}\right), \quad (2)$$

$$-\log \frac{I(\mathbf{r})}{I_0} = \int_{\mathbf{r}} \mu(\mathbf{x}) d\mathbf{l},$$

where $\mu(\mathbf{x})$ is the attenuation coefficient at spatial location $\mathbf{x}$ and $d\mathbf{l}$ is the differential path length along the ray. In other words, an X-ray pixel is not a sum of surface colors; it is the integral of attenuation contributions from different materials along a single path. What matters is how much of each material occupies the path, rather than how strong the local response is at one boundary point.

We represent the scene as a Gaussian set $\mathcal{G} = \{g_i\}_{i=1}^N$, where each Gaussian is written as

$$g_i = (\mathbf{p}_i, \Sigma_i, \rho_i, \mathbf{c}_i), \quad (3)$$

denoting the center, covariance, density, and appearance/radiative parameters, respectively. For any view π_v , the R2-Gaussian-style radiative renderer can be written as

$$\hat{I}_v = \mathcal{R}(\mathcal{G}; \pi_v). \quad (4)$$

Except for our new SPS, GAP, and ADM modules, the radiative rasterization, densification, and differentiable voxelization settings follow R2-Gaussian [ZLC*24]. Importantly, X-Gaussian and R2-Gaussian have already aligned the renderer with X-ray radiative projection. The mismatch we target lies mainly in the Gaussian evolution rule rather than in the imaging equation itself. In this baseline, Gaussians are still usually duplicated or split according to projection-error gradients. This is effective for improving surface coverage in natural scenes, but in CT it tends to keep pushing the evolving Gaussian set toward high-gradient anatomical boundaries. The model therefore uses the right X-ray renderer, yet still organizes Gaussians with a visible-light-style growth logic, which is precisely the issue our three modules address.

3.3. Overall Framework

As shown in Figure 2, **XRA-GS** instantiates the attenuation-aligned Gaussian evolution pipeline as three sequential modules embedded into the explicit radiative Gaussian backbone of R2-Gaussian. SPS seeds Gaussians into more relevant anatomical regions during initialization using the coarse FDK reconstruction. GAP then periodically reclaims locally redundant Gaussians after densification, preventing high-gradient boundaries from dominating Gaussian

growth. Finally, ADM refines the density distribution of retained Gaussians based on position-aware context during later optimization. In short, we do not redesign X-ray physics; we reorganize the optimization loop itself into path anchoring, boundary de-redundancy, and local refinement. Algorithm 1 summarizes the complete training procedure.

3.4. Support-Profile Seeding (SPS)

SPS initializes Gaussians in anatomically relevant regions using support-profile cues from the coarse FDK reconstruction. Initialization strongly influences the downstream optimization trajectory, yet neither of the standard 3DGS initializers fits CT well: uniform sampling cannot distinguish the anatomical importance of air, soft tissue, and bone, while SfM-based point-cloud initialization is unstable under gray-value overlap and weak local texture in X-ray projection [CLW*24]. Crucially, we do not use a vague spatial prior. Instead, we draw on three concrete signals contained in the FDK reconstruction: foreground anatomical support, low-frequency attenuation layout, and a coarse intensity hierarchy that differentiates air, soft tissue, and bony structures. The sparse-view CT literature also widely treats FDK/FBP reconstruction as a coarse CT volume with streak artifacts whose main value lies in global structure layout and useful frequency content rather than reliable fine detail [WCS*22; LMC*23; LFZ*25; ZLC*24]. Based on these structural properties of the FDK reconstruction, we use the coarse density volume V_{FDK} produced by FDK and define a density-weighted sampling distribution over a discrete voxel grid Ω :

$$p(\mathbf{x}) = \frac{(\max(V_{\text{FDK}}(\mathbf{x}), \epsilon))^\gamma}{\sum_{\mathbf{x}' \in \Omega} (\max(V_{\text{FDK}}(\mathbf{x}'), \epsilon))^\gamma}, \quad (5)$$

where ϵ avoids zero probability in background regions and γ controls sampling concentration.

Algorithm 1 XRA-GS Training Pipeline

Require: Sparse projections $\mathcal{P}$, FDK volume V_{FDK} , iterations T

Ensure: Optimized Gaussian set $\mathcal{G}^*$

1: **Stage 1 — Initialization (SPS):**

2: Compute density-weighted sampling $p(\mathbf{x})$ from V_{FDK} // Eq. (5)

3: Mix with uniform: $q(\mathbf{x}) \leftarrow \alpha/|\Omega| + (1-\alpha)p(\mathbf{x})$ // Eq. (6)

4: Sample M centers; set scales and densities from V_{FDK}

5: **Stage 2 — Main optimization loop:**

6: **for** $t = 1, \dots, T$ **do**

7: Render $\hat{I}_v \leftarrow \mathcal{R}(\mathcal{G}; \pi_v)$ for observed views // Eq. (4)

8: Compute loss $\mathcal{L}$ // Eq. (16)

9: Update $\mathcal{G}$, ADM parameters via back-propagation

10: **if** $t \bmod T_{\text{densify}} = 0$ **then**

11: Standard densification (clone / split)

12: **if** $t \in [T_{\text{gap}}, T_{\text{gap}, \text{end}}]$ **then**

13: **GAP:** compute proximity s_i ; prune by mask m_i // Eqs. (7)–(8); $T_{\text{gap}} = 2K$, $T_{\text{gap}, \text{end}} = 20K$

14: **end if**

15: **end if**

16: **if** $t \geq T_{\text{adm}}$ **then**

17: **ADM:** predict $\tilde{o}(\mathbf{x})$; modulate ρ_{final} // Eq. (15)

18: **end if**

19: **end for**

20: **return** $\mathcal{G}^* \leftarrow \mathcal{G}$

Density weighting alone would over-concentrate Gaussians near the brightest bony structures and could mistake the strongest sparse-view FDK responses for precise geometry. We therefore mix uni-

Placeholder: `assets/fig/method_fig_pipeline.png`

Figure 2: Overall framework of XRA-GS. Given sparse X-ray projections $\{(I_v, \pi_v)\}_{v=1}^V$ and a coarse FDK reconstruction V_{FDK} , SPS (purple) performs path-anchored initialization along the attenuation support outside the optimization loop. An iterative training loop then interleaves ADM (orange, position-aware density modulation) and GAP (teal, gradient-aware redundancy recovery), supervised end-to-end by the rendering loss. Solid arrows denote data flow; dashed arrows denote gradient flow. The output is a novel-view X-ray projection together with the recovered tomographic volume.

form sampling with the support-profile distribution:

$$q(\mathbf{x}) = \alpha \frac{1}{|\Omega|} + (1 - \alpha)p(\mathbf{x}), \quad \mathbf{x} \in \Omega, \quad (6)$$

and independently sample M initial centers $\{\mathbf{x}_m\}_{m=1}^M$ from $q(\mathbf{x})$, where M is the number of initial Gaussians. This design preserves global coverage while allocating more Gaussians to anatomical regions with meaningful attenuation support, without collapsing entirely onto dense bony boundaries. After sampling the locations, we initialize Gaussian densities from the FDK values and set scale parameters according to the local sampling density, yielding 50K initial Gaussians. In all experiments, α and γ are selected by a small local search and then fixed to avoid scene-specific manual tuning.

3.5. Geometry-aware Pruning (GAP)

GAP periodically removes locally redundant Gaussians based on spatial proximity and gradient activity. In CT, gradient-driven densification often

keeps duplicating Gaussians around bone-soft tissue boundaries because these regions produce larger projection-error gradients. However, high gradient does not always imply missing coverage; it may also indicate that the same anatomical boundary is being represented repeatedly by neighboring Gaussians. Instead of following the FSGS logic of adding support points in sparse regions [ZFJW24], GAP explicitly measures local crowding in world coordinates and uses it for redundancy recovery in over-densified regions.

Specifically, for each Gaussian G_i , let $\mathcal{N}_K(i)$ denote the index set of its K nearest neighbors excluding itself. We compute the mean Euclidean distance to these neighbors as a proximity score:

$$s_i = \frac{1}{K} \sum_{j \in \mathcal{N}_K(i)} \|\mathbf{p}_i - \mathbf{p}_j\|_2. \quad (7)$$

When s_i falls below a threshold τ , the Gaussian lies in an overly crowded region. To avoid deleting Gaussians that are still learning useful structure, we fur-




Figure 3: Support-Profile Seeding (SPS). (a) From the sparse projections, a coarse FDK volume is reconstructed; the horizontal attenuation profile along an anatomical scan line exposes the support-profile cue (purple curve). (b) SPS converts the attenuation magnitude into a density-weighted sampling map (viridis) and mixes it with a uniform prior, so capacity concentrates on anatomically relevant support while a small uniform term retains global coverage. (c) Compared with random initialization (gray dots), SPS produces an initial Gaussian set (purple dots) already aligned with the attenuation support before any optimization step.

ther introduce a gradient-activity filter. Let $\bar{g}_i$ denote the accumulated L_1 statistic of the positional gradient of this Gaussian over the current densification cycle. We define a retention mask $m_i \in \{0, 1\}$ as

$$m_i = \begin{cases} 1, & s_i \geq \tau \vee \bar{g}_i \geq \delta, \\ 0, & \text{otherwise.} \end{cases} \quad (8)$$

where $m_i = 0$ means that the Gaussian is treated as a redundancy candidate in the current cycle. In implementation, GAP does not replace the original 3DGS densification; it is a periodic recovery step applied afterwards. We first generate candidate Gaussians through the standard duplication/splitting rule and then use m_i to filter redundant ones. In each cycle, we prune at most a fraction β_{prune} of the total




Figure 4: Geometry-aware Pruning (GAP). (a) After standard gradient-driven densification, Gaussians accumulate around high-attenuation boundaries, producing visible over-densification clusters. (b) For each Gaussian, GAP measures local crowding by the mean Euclidean distance to its K nearest neighbors and combines it with the accumulated gradient activity over the current densification cycle: a Gaussian is flagged as a redundancy candidate (red $\times$) only if both KNN distance is below τ and gradient activity is below δ ; otherwise it is retained (green $\circ$). (c) After removing at most a fraction β_{prune} of the redundant candidates, the cleaned distribution preserves the attenuation support while freeing capacity for under-served interior regions.

Gaussians, and the module is triggered periodically from 2K to 20K iterations. To avoid suddenly making local support too sparse after pruning, we only shrink the covariance of newly retained Gaussians:

$$\Sigma_i^{\text{new}} = \eta \Sigma_i, \quad 0 < \eta < 1. \quad (9)$$

GAP therefore regularizes the densification process by suppressing locally redundant Gaussians, preventing high-gradient boundaries from continuously monopolizing representation capacity.

3.6. Adaptive Density Modulation (ADM)

ADM predicts position-dependent density offsets to refine the density distribution of retained Gaussians. Even after better initialization and redundancy control, density updates remain highly heterogeneous across anatomy: bone usually corresponds to higher attenuation, air regions in the lung should remain low-density, and soft tissue often requires smoother transitions. If all Gaussians share the same point-wise gradient update rule, the model struggles to express this spatial heterogeneity. We therefore encode Gaussian centers with K-Planes and use a lightweight decoder that predicts only density-related offsets and confidence gates, without rewriting Gaussian centers, covariances, or other attributes. The role of ADM is local density refinement rather than serving as the sole source of performance.

It is important to emphasize that, although ADM contains an MLP, it is not an offline pre-trained prediction head and does not rely on extra density supervision. Both the K-Planes features and the dual-head MLP parameters are updated only during optimization on the current case, and they are trained end-to-end through the main rendering loss. In this sense, ADM is an online joint-optimization branch that adaptively amplifies or suppresses density responses in different regions while the main radiative Gaussian backbone is being trained.

Given a position $\mathbf{x}$, we first project it onto normalized 2D coordinates $\Pi_{uv}(\mathbf{x})$ on three orthogonal planes and then bilinearly sample the corresponding feature maps:

$$f_{uv}(\mathbf{x}) = \mathcal{B}(P_{uv}, \Pi_{uv}(\mathbf{x})), \quad uv \in \{xy, xz, yz\}, \quad (10)$$

$$F(\mathbf{x}) = \text{concat}(f_{xy}(\mathbf{x}), f_{xz}(\mathbf{x}), f_{yz}(\mathbf{x})), \quad (11)$$

where $\mathcal{B}(\cdot)$ denotes bilinear sampling, P_{xy} , P_{xz} , and P_{yz} are the three orthogonal feature planes, and

$\text{concat}(\cdot)$ denotes channel-wise concatenation. K-Planes encoding was originally introduced for explicit spatiotemporal radiance fields and provides continuous spatial context at low cost [FYT*23]. Based on this encoding, a dual-head MLP predicts a density offset $\Delta\sigma(\mathbf{x})$ and a confidence gate $g(\mathbf{x})$:

$$(\Delta\sigma(\mathbf{x}), g(\mathbf{x})) = h(F(\mathbf{x})). \quad (12)$$

We then center the offsets by the batch mean and apply a view-dependent scaling:

$$\tilde{o}(\mathbf{x}) = \beta(t) s_{\text{view}} r_{\text{max}} g(\mathbf{x}) \cdot (\Delta\sigma(\mathbf{x}) - \overline{\Delta\sigma}), \quad (13)$$

where $\overline{\Delta\sigma}$ is the mean density offset of the current batch, r_{max} is the maximum modulation amplitude, and s_{view} is a view-count scaling factor (concretely, $s_{\text{view}} = 0.5, 0.7, 1.0$ for 2, 3, and 4 input views, respectively). To prevent modulation from disrupting basic geometric convergence early in training, we use a warm-up/hold/decay schedule for $\beta(t)$:

$$\beta(t) = \begin{cases} \frac{t}{t_w}, & 0 \leq t < t_w, \\ 1, & t_w \leq t < t_h, \\ \exp\left(-\frac{t-t_h}{\kappa}\right), & t \geq t_h. \end{cases} \quad (14)$$

where t_w marks the end of the warm-up phase, t_h marks the end of the hold phase, and κ is the exponential decay time constant. The final density is

$$\rho_{\text{final}}(\mathbf{x}) = \rho_{\text{base}}(\mathbf{x}) \cdot (1 + \tilde{o}(\mathbf{x})). \quad (15)$$

Therefore, ADM learns not a global density boost, but a spatial-context-aware relative correction that better matches the locally continuous yet regionally heterogeneous structure of CT data.

3.7. Training Loss

The three modules of **XRA-GS** and the radiative rendering backbone of R2-Gaussian form a single end-to-end training pipeline. During training, we supervise reprojection results on the observed views

Placeholder: `assets/fig/method_fig_adm.png`

Figure 5: Adaptive Density Modulation (ADM). (a) Three orthogonal K-Planes feature maps (*viridis*) encode position-dependent spatial context for every Gaussian center. (b) The concatenated K-Planes feature is decoded by a lightweight dual-head MLP that predicts a density offset $\Delta\sigma(\mathbf{x})$ (orange) and a confidence gate $g(\mathbf{x}) \in [0, 1]$ (slate); the modulated density combines them through $\rho_{\text{final}} = \rho_{\text{base}} \cdot (1 + \tilde{o})$. (c) The resulting density difference map (*inferno*) shows that ADM amplifies density in regions where the attenuation support requires fine detail and damps modulation where evidence is uncertain.

with the following overall loss:

$$\mathcal{L} = \mathcal{L}_1 + \lambda_{\text{dssim}} \mathcal{L}_{\text{dssim}} + \lambda_{3\text{DTV}} \mathcal{L}_{3\text{DTV}} + \lambda_{\text{pTV}} \mathcal{L}_{\text{pTV}}, \quad (16)$$

where

$$\mathcal{L}_1 = \frac{1}{V} \sum_{v=1}^V \|I_v - \hat{I}_v\|_1, \quad (17)$$

$$\mathcal{L}_{\text{dssim}} = \frac{1}{V} \sum_{v=1}^V \text{DSSIM}(I_v, \hat{I}_v), \quad (18)$$

$$\mathcal{L}_{3\text{DTV}} = \sum_{i=1}^N \sum_{d \in \{x, y, z\}} |\nabla_d \rho_i|, \quad (19)$$

$$\mathcal{L}_{\text{pTV}} = \sum_{p \in \{xy, xz, yz\}} \omega_p(t) \text{TV}(P_p). \quad (20)$$

Here, $\omega_p(t)$ is the time-varying regularization weight for the p -th feature plane; it follows the same warm-

up/hold/decay schedule as $\beta(t)$, ramping linearly from 0 to ω_0 over the first t_w iterations, holding at ω_0 until t_h , then decaying exponentially with time constant κ . $\mathcal{L}_1$ and $\mathcal{L}_{\text{dssim}}$ jointly constrain projection reconstruction quality, $\mathcal{L}_{3\text{DTV}}$ encourages spatial smoothness in volumetric density, and $\mathcal{L}_{\text{pTV}}$ imposes planar total-variation regularization on the K-Planes features. During training, the plane features and MLP parameters of ADM are jointly updated with the backbone Gaussian parameters under the same loss, so density modulation is learned online rather than fixed in advance. Since SPS only changes initialization, GAP is triggered periodically during training, and ADM applies lightweight density modulation, the basic X-ray rendering process at inference remains unchanged and the extra cost is concentrated in training.

4. Experiments

4.1. Experimental Setup

Data and task setting. We evaluate on the five-organ dataset from the X-Gaussian repository [CLW*24]: Chest (LIDC-IDRI) and Head, Abdomen, Foot, and Pancreas (open scientific visualization dataset). Following the NAF/X-Gaussian preprocessing and projection protocol, we construct sparse-view settings by retaining only 2, 3, or 4 input projections per scene. Models are trained exclusively on observed projections and evaluated on held-out novel angles.

Baselines. We compare against six methods in two groups. *Direct X-ray/CT baselines:* X-Gaussian [CLW*24] (radiative Gaussian splatting), R2-Gaussian [ZLC*24] (reconstruction-consistency regularization), and X-Field [WTW*25] (material-aware implicit modeling) represent the state of the art for X-ray novel view synthesis. *Natural-scene sparse-view 3DGS references:* CoR-GS [ZLY*24],

DNGaussian [LZB*24], and FSGS [ZFJW24] characterize how gradient-driven densification degrades under tomographic projection. X-Field is re-evaluated under our identical protocol and denoted as X-Field* in the result tables and figures.

Metrics. Following the official X-Gaussian evaluation script, PSNR2D and SSIM2D are computed on $[0, 1]$ -normalized projections averaged over all held-out views of each scene and then over the five organs. PSNR2D is $20 \log_{10}(1/\sqrt{\text{MSE}})$; SSIM2D uses an 11×11 Gaussian window ($\sigma = 1.5$, $C_1 = 0.01^2$, $C_2 = 0.03^2$).

Implementation details. SPS uses $\alpha = 0.2$, $\gamma = 1.0$, and 50K initial Gaussians. GAP uses $K = 5$ neighbors, threshold $\tau = 0.015$, maximum pruning ratio $\beta_{\text{prune}} = 0.02$, and gradient threshold $\delta = 0.0002$. ADM uses K-Planes at resolution 64×32 and activates after 15K iterations. All models are trained for 30K iterations.

4.2. Main Results

Table 1 reports PSNR2D across all five organs and three sparsity levels. **XRA-GS** achieves the best five-organ average at every sparsity level. Gains over R2-Gaussian are $+0.11/ + 0.39/ + 0.07$ dB under 2/3/4-view settings; gains over X-Field* are $+0.87/ + 3.25/ + 2.79$ dB. The most pronounced margin occurs at 3 views, consistent with our central argument: capacity misallocation is most directly addressable when projection evidence is sufficient to guide redistribution yet too limited for standard densification to converge to a spatially consistent solution.

Table 2 reports SSIM2D averages for the three methods with accessible logs under our unified protocol. **XRA-GS** leads at 2- and 3-view settings and ties R2-Gaussian at the third decimal place in 4-view, confirming the same trend as PSNR2D.

Figure 6 provides qualitative comparisons across

all five organs under the 2-view setting. The natural-scene sparse-view baselines (CoR-GS, DNGaussian, FSGS) produce characteristic boundary leakage and streak artifacts under tomographic projection; the direct X-ray baselines (X-Gaussian, R2-Gaussian, X-Field*) preserve global attenuation structure more faithfully. **XRA-GS** further reduces background diffusion and local over-accumulation near high-contrast anatomical boundaries, consistent with the quantitative gains in Table 1.

Figure 7 provides two complementary analyses. The upper portion enlarges the Chest and Foot regions marked in Figure 6, isolating boundary and structure recovery at the local scale. The lower portion evaluates a Chest 2-view case at three novel angles (0° , 45° , 90°), confirming that the advantage of **XRA-GS** over R2-Gaussian is stable across viewing directions rather than concentrated at a single angle.

Same-protocol comparison with X-Field. The X-Field* row in Table 1 confirms that the gains of **XRA-GS** are not an artifact of protocol asymmetry. Re-evaluated under our identical 2/3/4-view setting, X-Field* trails **XRA-GS** by a consistent average margin across all sparsity levels.

4.3. Module and Initialization Analysis

Table 3 reports a progressive ablation starting from the R2-Gaussian backbone. SPS raises the 2-view average PSNR2D from 21.27 to 21.44 dB, establishing that path-anchored initialization provides a meaningful head start in the most under-determined regime. GAP delivers the largest PSNR2D increment, improving the 3-view average from 28.01 to 28.22 dB by reclaiming capacity from over-densified boundary regions. The final **XRA-GS** row matches or improves the available PSNR2D measurements across all view counts, while intermediate SSIM2D values are pending from the experiment logs.

Table 1: Main PSNR2D comparison under the unified five-organ, 2/3/4-view protocol. Blank entries denote same-protocol results that are unavailable in the source logs, so we do not estimate or extrapolate them. Best and second-best results are highlighted in bold and underline, respectively.

Method	Chest			Head			Abdomen		
	2v	3v	4v	2v	3v	4v	2v	3v	4v
DNGaussian	14.30	16.02	15.76	15.72	18.04	16.24	19.33	23.51	24.72
CoR-GS	14.96	18.55	20.29	16.24	20.36	23.40	18.50	22.45	23.03
FSGS	17.86	19.99	21.03	19.43	21.53	24.91	18.75	23.34	23.63
X-Gaussian	17.79	20.29	20.90	19.62	21.52	24.46	18.93	23.54	23.85
R2-Gaussian	<u>21.05</u>	<u>26.22</u>	<u>25.47</u>	<u>23.54</u>	<u>26.60</u>	28.66	<u>24.85</u>	<u>29.26</u>	<u>30.56</u>
X-Field*									
XRA-GS	21.26	26.81	25.83	23.63	26.63	<u>28.52</u>	24.97	29.73	30.89

Method	Foot			Pancreas			Avg		
	2v	3v	4v	2v	3v	4v	2v	3v	4v
DNGaussian	20.17	23.51	19.33	16.15	22.43	25.47	17.13	20.70	20.30
CoR-GS	20.17	25.53	26.83	16.15	23.25	23.96	17.20	22.03	23.50
FSGS	<u>22.17</u>	25.37	27.08	<u>23.03</u>	25.27	26.22	20.25	23.10	24.57
X-Gaussian	22.34	25.48	26.92	23.21	25.12	25.55	20.38	23.19	24.34
R2-Gaussian	19.39	<u>28.48</u>	30.04	17.83	<u>28.61</u>	30.90	<u>21.33</u>	<u>27.83</u>	<u>29.13</u>
X-Field*							20.57	24.97	26.41
XRA-GS	19.55	28.49	<u>29.86</u>	17.80	29.43	30.90	21.44	28.22	29.20

Figure 8 translates the progressive ablation into local visual evidence. Residual maps $|\hat{I} - I_{gt}|$ are shown below each prediction. This visual analysis is used to inspect where the progressive modules change the reconstruction, while the corresponding intermediate SSIM2D values remain pending from the experiment logs.

Figure 9 is reserved for visualizing the cumulative effect of the three modules on Gaussian spatial

distribution once the corresponding Gaussian-center and count logs are available. Until those logs are confirmed, the mechanism is evaluated through the available quantitative ablation in Table 3.

Table 4 (left) isolates the initialization strategy. SPS-guided seeding outperforms random initialization at all sparsity levels; the full **XRA-GS** further improves upon SPS-only through GAP and ADM,




Figure 6: Qualitative comparison on five organs, 2-view setting. Rows: Chest, Head, Abdomen, Foot, Pancreas (top to bottom). Columns: GT, CoR-GS, DNGaussian, FSGS, X-Field*, X-Gaussian, R2-Gaussian, and **XRA-GS** (Ours, last column). Natural-scene references produce streak artifacts and boundary leakage; direct X-ray baselines preserve global attenuation structure; **XRA-GS** further suppresses local over-accumulation and background diffusion. Rightmost column: error maps (inferno, shared vmax). Red boxes mark regions enlarged in Figure 7.

Table 2: SSIM2D comparison (five-organ average). Per-organ SSIM2D logs for CoR-GS, DNGaussian, FSGS, and X-Gaussian are unavailable under our unified 2/3/4-view protocol; we therefore report same-protocol averages only for the three methods with accessible logs. **Bold:** best; underline: second-best.

Method	2v Avg.	3v Avg.	4v Avg.
X-Field*	0.717	0.850	0.880
R2-Gaussian [ZLC*24]	<u>0.794</u>	<u>0.903</u>	0.924
XRA-GS (Ours)	0.797	0.904	0.924

Table 3: Progressive component ablation (five-organ average). Modules are added incrementally on top of the R2-Gaussian backbone. GAP delivers the largest available PSNR2D increment, while SPS benefits the most under-determined 2-view regime. Intermediate SSIM2D values are pending from experiment logs, so the current table supports PSNR2D trends only for intermediate configurations. **Bold:** best per column.

Configuration	PSNR2D $\uparrow$ (dB)			SSIM2D $\uparrow$		
	2v	3v	4v	2v	3v	4v
Baseline (R2-Gaussian)	21.27	27.80	29.10	–	–	–
+ SPS	21.44	28.01	29.16	–	–	–
+ SPS + GAP	21.44	28.22	29.20	–	–	–
XRA-GS (SPS + GAP + ADM)	21.44	28.22	29.20	0.797	0.904	0.924

SSIM2D for intermediate configurations is pending from experiment logs.

confirming that path-anchored initialization and capacity recovery are complementary operations.

4.4. View-count Analysis

Table 4 (right) reports the view-count trend for four representative methods. **XRA-GS** achieves the

Placeholder: `assets/fig/experiment_fig_qual_detail_consistency.png`
 Upper: zoom-in Chest + Foot (4 cols: GT / R2-Gaussian/ X-Gaussian / **XRA-GS**)
 Lower: Chest 2-view at $0^\circ/45^\circ/90^\circ$ (3 rows $\times$ 4 cols)

Figure 7: Local detail and cross-view consistency. Upper half: Zoom-ins from the red boxes in Figure 6 for Chest (left) and Foot (right), comparing GT, R2-Gaussian, X-Gaussian, and **XRA-GS**. **XRA-GS** reduces streak spread and boundary diffusion in both regions. Lower half: A Chest 2-view case rendered at 0° , 45° , and 90° . The quality advantage of **XRA-GS** persists across all three angles, indicating consistent global attenuation reconstruction rather than angle-specific improvement.

Placeholder: `assets/fig/experiment_fig_ablation_visual.png`
 Chest 3-view — top: GT / Baseline / +SPS / +SPS+GAP / Full **XRA-GS**
 bottom: residual maps $|\hat{I} - I_{gt}|$ (inferno); ΔSSIM2D annotated under each col

Figure 8: Progressive module ablation, Chest 3-view. Top row: rendered projections for GT, Baseline (R2-Gaussian), +SPS, +SPS+GAP, and full **XRA-GS**. Bottom row: residual maps $|\hat{I} - I_{gt}|$ (inferno, shared $v_{\max}$ with Figure 6). Quantitative intermediate SSIM2D values are pending from experiment logs and are therefore not used for module-level claims here.

highest five-organ average PSNR2D at every sparsity level. The relative gain over R2-Gaussian is largest at 3 views, where capacity misallocation is most directly observable; at 4 views the reconstruction-consistency

regularization in R2-Gaussian partially closes the gap, consistent with the applicability boundary discussed in Section 4.7.




Figure 9: Gaussian spatial distribution evolution. Four panels (left to right): uniform initialization, after SPS, after GAP pruning, and full **XRA-GS**, for a representative Chest organ. Top strip: FDK attenuation profile (magma shading). The final figure will use logged Gaussian centers and counts rather than schematic estimates.

Table 4: Left: Initialization strategy ablation (five-organ average PSNR2D, dB). SPS-guided seeding outperforms random initialization at all sparsity levels; the full **XRA-GS** further gains from GAP and ADM, confirming that path-anchored initialization and capacity recovery are complementary. **Right: View-count trend** for four representative methods (five-organ average PSNR2D, dB). **XRA-GS** leads at every sparsity level; the relative advantage over R2-Gaussian is largest at 3 views, where capacity misallocation is most directly observable. **Bold:** best; underline: second-best.

Initialization	2v	3v	4v
Random (no SPS)	21.27	27.80	29.10
SPS-guided (no GAP, ADM)	<u>21.44</u>	<u>28.01</u>	<u>29.16</u>
XRA-GS (full)	21.44	28.22	29.20

Five-organ avg PSNR2D $\uparrow$ (dB). Per-organ and SSIM2D data in supplementary.

Method	2v	3v	4v
FSGS [ZFJW24]	20.25	23.10	24.57
X-Gaussian [CLW*24]	20.38	23.19	24.34
R2-Gaussian [ZLC*24]	<u>21.33</u>	<u>27.83</u>	<u>29.13</u>
XRA-GS (Ours)	21.44	28.22	29.20

Five-organ avg PSNR2D $\uparrow$ (dB). CoR-GS and DNGaussian omitted for brevity (full data in Table 1).

4.5. Computational Cost

Table 5 is reserved for training and inference cost under the 3-view setting. The current manuscript keeps the table structure in place, but Time, final Gaussian count, GPU memory, and FPS must be filled

from experiment logs before making quantitative efficiency claims.

4.6. Hyperparameter Sensitivity

Figure 10 is reserved for local sweeps of four key hyperparameters in the Chest 3-view setting: the SPS mixture coefficient α , the GAP proximity threshold

Table 5: Computational cost comparison under the 3-view setting (five-organ average). Train Time: end-to-end training minutes. #Gaussians: final primitive count ($\times 10^3$) after densification/pruning. GPU Mem: peak memory during training (GB). FPS: rendering frames per second at inference. PSNR2D: five-organ average (dB), from Table 1. Cost columns will be filled from matched experiment logs before quantitative efficiency claims are made. **Bold:** best per column.

Method	Time↓ (min)	#Gauss.↓ (K)	Mem.↓ (GB)	FPS↑	PSNR2D↑ (dB)
X-Gaussian [CLW*24]	–	–	–	–	23.19
R2-Gaussian [ZLC*24]	–	–	–	–	27.83
XRA-GS (Ours)	–	–	–	–	28.22

“–” indicates values pending from experiment logs requested in `assets/ask/experiment_2026-05-23_T01_T02_T04_main_metrics_efficiency.md`. PSNR2D is taken from Table 1.

τ , the GAP pruning ratio β_{prune} , and the ADM activation iteration. The default configuration ($\alpha = 0.2$, $\tau = 0.015$, $\beta_{\text{prune}} = 2\%$, ADM activation at 15K) will be assessed once the sweep metrics are available from the experiment logs.

Placeholder: `assets/fig/experiment_fig_hparam.png`
 2×2 panels: (a) SPS α ; (b) GAP τ ; (c) GAP β_{prune} ; (d) ADM activation
 Primary y-axis: SSIM2D (solid Crimson); secondary: PSNR2D (dashed); • = default

Figure 10: Hyperparameter sensitivity, Chest 3-view. Each panel sweeps one hyperparameter with all others fixed at the default (marked •). Solid Crimson: SSIM2D↑; dashed: PSNR2D↑. (a) SPS mixture coeff. $\alpha \in \{0.0, 0.25, 0.5, 0.75, 1.0\}$; (b) GAP threshold τ ; (c) GAP pruning ratio $\beta_{\text{prune}} \in \{0, 1, 2, 5, 10\}\%$; (d) ADM activation iteration (fraction of 30K iters). The final sweep figure will be populated from experiment logs.

4.7. Discussion and Limitations

The experiments confirm that in the moderately sparse regime, the primary bottleneck is capacity allocation rather than Gaussian count. Gains are most pronounced at 3 views, where sufficient projection evidence exists for path-anchored initialization and boundary de-redundancy to take effect, yet too little evidence for standard densification to converge to a spatially consistent solution.

Two boundaries constrain the current method. At 2 views the FDK reconstruction used by SPS is insufficiently reliable, so the initialization advantage partially erodes. At 4 views the reconstruction-consistency regularization in R2-Gaussian is strong enough to compensate for boundary-chasing effects, reducing the relative advantage of **XRA-GS**. Figure 11 illustrates both failure modes: in Foot 2-view

and Pancreas 2-view, fine-grained high-frequency detail remains unrecoverable when projection evidence is fundamentally insufficient, regardless of initialization or capacity management.

Several directions remain open. First, SPS relies on FDK quality, which degrades sharply below two views; learned or uncertainty-aware initializers could relax this constraint. Second, the GAP threshold and pruning ratio are fixed organ-agnostically; organ-adaptive scheduling may improve per-organ performance. Third, all experiments assume parallel-beam geometry; extension to cone-beam and dynamic CT acquisition is left for future work. Additional per-organ SSIM2D breakdowns and single-module ablations will be included only after the corresponding same-protocol logs are confirmed.

5. Conclusion

XRA-GS demonstrates that visible-light densification rules impose a systematic capacity-allocation failure in transmissive X-ray imaging that persists even after the rendering equation is physically corrected. By intervening at three stages of Gaussian evolution—path-anchored initialization (SPS), boundary redundancy recovery (GAP), and position-dependent local refinement (ADM)—the framework redirects representation capacity from overdensified anatomical boundaries toward the interior attenuation-relevant structures that conventional gradient-driven growth neglects. Across a five-organ benchmark with 2/3/4 input views against six representative baselines, **XRA-GS** achieves the best average PSNR2D at every sparsity level, with SSIM2D reported for methods whose same-protocol logs are available. The strongest PSNR2D improvement appears in the moderately sparse 3-view regime, where sufficient projection evidence allows capacity reallocation to take effect.

References

- [AK84] ANDERSEN, A. H. and KAK, A. C. “Simultaneous Algebraic Reconstruction Technique (SART): A Superior Implementation of the ART Algorithm”. *Ultrasonic Imaging* 6.1 (1984), 81–94 1, 5.
- [CLTS24] CHARATAN, DAVID, LI, SIZHE LESTER, TAGLIASACCHI, ANDREA, and SITZMANN, VINCENT. “pixelSplat: 3D Gaussian Splats from Image Pairs for Scalable Generalizable 3D Reconstruction”. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2024, 19457–19467 4.
- [CLW*24] CAI, YUANHAO, LIANG, YIXUN, WANG, JIAHAO, et al. “Radiative Gaussian Splatting for Efficient X-ray Novel View Synthesis”. *European Conference on Computer Vision (ECCV)*. 2024 2, 5, 7, 11, 16, 17.
- [CWY*24] CAI, YUANHAO, WANG, JIAHAO, YUILLE, ALAN, et al. “Structure-Aware Sparse-View X-ray 3D Reconstruction”. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2024, 11174–11183 5.
- [CXZ*24] CHEN, YUEDONG, XU, HAOFEI, ZHENG, CHUANXIA, et al. “MVSplat: Efficient 3D Gaussian Splatting from Sparse Multi-View Images”. *European Conference on Computer Vision (ECCV)*. 2024 4.
- [DLZR22] DENG, KANGLE, LIU, ANDREW, ZHU, JUN-YAN, and RAMANAN, DEVA. “Depth-Supervised NeRF: Fewer Views and Faster Training for Free”. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2022 4.
- [FDK84] FELDKAMP, L. A., DAVIS, L. C., and KRESS, J. W. “Practical Cone-Beam Algorithm”. *Journal of the Optical Society of America A* 1.6 (1984), 612–619 1, 3, 5.
- [FYT*23] FRIDOVICH-KEIL, SARA, YU, GIANG, TANCIK, MATTHEW, et al. “K-Planes: Explicit Radiance Fields in Space, Time, and Appearance”. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2023 4, 10.
- [GL24] GUÉDON, ANTOINE and LEPETIT, VINCENT. “SuGaR: Surface-Aligned Gaussian Splatting for Efficient 3D Mesh Reconstruction and High-Quality Mesh Rendering”. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2024 2, 3.
- [GWS*24] GUAN, QIANG, WANG, JIARUI, SHEN, LIYUE, et al. “NAF: Neural Attenuation Fields for Sparse-View X-ray Imaging”. *Medical Image Analysis* (2024) 5.
- [HYC*24] HUANG, BINBIN, YU, ZHAO, CHEN, ANPEI, et al. “2D Gaussian Splatting for Geometrically Accurate Radiance Fields”. *ACM SIGGRAPH 2024 Conference Papers*. 2024 2, 3.




Figure 11: Representative failure cases, 2-view setting. Rows: Foot (top), Pancreas (bottom). Columns: GT, R2-Gaussian, **XRA-GS**, and **XRA-GS** error map (inferno, shared vmax with Figure 6). Zoom-in boxes mark the primary failure region. **XRA-GS** reduces overall error spread relative to R2-Gaussian but cannot resolve fine structural detail when projection evidence is fundamentally absent. Failure descriptions (≤ 25 words) are inset at the right of each row.

- [JTA21] JAIN, AJAY, TANCIK, MATTHEW, and ABBEEL, PIETER. “Putting NeRF on a Diet: Semantically Consistent Few-Shot View Synthesis”. *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*. 2021 4.
- [KKLD23] KERBL, BERNHARD, KOPANAS, GEORGIOS, LEIMKÜHLER, THOMAS, and DRETTAKIS, GEORGE. “3D Gaussian Splatting for Real-Time Radiance Field Rendering”. *ACM Transactions on Graphics* 42.4 (2023), 1–14 2, 3.
- [KS01] KAK, AVINASH C. and SLANEY, MALCOLM. *Principles of Computerized Tomographic Imaging*. Philadelphia: Society for Industrial and Applied Mathematics, 2001 2.
- [Kua*25] KUANG, YUANHAO et al. “Layer-Based 3D Gaussian Splatting for Sparse-View CT Reconstruction”. *Proceedings of the IEEE/CVF International Conference on Computer Vision*. 2025 5.
- [LFL*25] LIU, ZHENTAO, FANG, YU, LI, CHANGJIAN, et al. “Geometry-Aware Attenuation Learning for Sparse-View CBCT Reconstruction”. *IEEE Transactions on Medical Imaging* 44.2 (2025), 1083–1097 5.
- [LFY*25] LIU, YIFAN, FAN, KEYU, YU, WEIHAO, et al. “MonoSplat: Generalizable 3D Gaussian Splatting from Monocular Depth Foundation Models”. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2025 4.
- [LFZ*25] LI, YINGTAI, FU, XUEMING, ZHAO, SHANG, et al. “3DGR-CT: Sparse-view CT Reconstruction with a 3D Gaussian Representation”. *Medical Image Analysis* 103 (2025), 103585 7.
- [LMC*23] LI, ZILONG, MA, CHENGLONG, CHEN, JIE, et al. “Learning to Distill Global Representation for Sparse-View CT”. *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*. 2023, 21196–21207 1, 7.
- [LYX*24] LU, TAO, YU, MULIN, XU, LINNING, et al. “Scaffold-GS: Structured 3D Gaussians for View-Adaptive Rendering”. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2024 3.
- [LZB*24] LI, JIAHE, ZHANG, JIAWEI, BAI, XIAO, et al. “DNGaussian: Optimizing Sparse-View 3D Gaussian Radiance Fields with Global-Local Depth Normalization”. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2024 2, 4, 12.
- [Max95] MAX, NELSON. “Optical Models for Direct Volume Rendering”. *IEEE Transactions on Visualization and Computer Graphics* 1.2 (1995), 99–108 2.
- [MST*21] MILDENHALL, BEN, SRINIVASAN, PRATUL P., TANCIK, MATTHEW, et al. “NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis”. *Communications of the ACM* 65.1 (2021), 99–106 4.
- [NBM*22] NIEMEYER, MICHAEL, BARRON, JONATHAN T., MILDENHALL, BEN, et al. “RegNeRF: Regularizing Neural Radiance Fields for View Synthesis from Sparse Inputs”. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2022 4.

- [NMKW24] NAVANEET, K., MEIBODI, KOOSHK, KOIKE-AKINO, TOSHIKI, and WANG, YIBO. “Compact 3D Gaussian Splatting for Sparse View Rendering”. *arXiv preprint arXiv:2311.18159* (2024) 3.
- [PRK25] PARK, HYUNWOO, RYU, GUN, and KIM, WONJUN. “DropGaussian: Structural Regularization for Sparse-view Gaussian Splatting”. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2025, 21600–21609 4.
- [RNSZ24] RÖSSLE, BARBARA, NIESSNER, MATTHIAS, SEIDEL, HANS-PETER, and ZOLLHÖFER, MICHAEL. “Densification Strategies for 3D Gaussian Splatting”. *arXiv preprint arXiv:2401.17151* (2024) 2.
- [TRZ*24] TANG, JIAXIANG, REN, JIAWEI, ZHOU, HANG, et al. “DreamGaussian: Generative Gaussian Splatting for Efficient 3D Content Creation”. *International Conference on Learning Representations (ICLR)*. 2024 3.
- [WCLL23] WANG, GUANGCONG, CHEN, ZHAOXI, LOY, CHEN CHANGE, and LIU, ZIWEI. “SparseNeRF: Distilling Depth Ranking for Few-Shot Novel View Synthesis”. *arXiv preprint arXiv:2303.16196* (2023) 4.
- [WCS*22] WANG, YANLI, CHAO, LIANYING, SHAN, WENQI, et al. “Improving the Quality of Sparse-view Cone-Beam Computed Tomography via Reconstruction-Friendly Interpolation Network”. *Proceedings of the Asian Conference on Computer Vision (ACCV)*. 2022, 2255–2269 7.
- [WHL*24] WANG, ZHE, HUANG, YUKUN, LIANG, YUXIANG, et al. “Reducing Image Artifacts in Sparse Projection CT Using Conditional Generative Adversarial Networks”. *Scientific Reports* 14 (2024), 3889 1.
- [WLG*25] WU, SHAOKAI, LU, YUXIANG, GUO, YAPAN, et al. “Discretized Gaussian Representation for Tomographic Reconstruction”. *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*. 2025, 25073–25082 2, 5.
- [WN19] WILLEMINK, MICHIEL J. and NOËL, PETER B. “The Evolution of Image Reconstruction from Filtered Back Projection to Artificial Intelligence: A Review”. *European Radiology* 29 (2019), 4805–4813 2, 5.
- [WRW*24] WU, ZIYI, REN, JIAWEI, WANG, XIAO, et al. “4D Gaussian Splatting for Real-Time Dynamic Scene Rendering”. *arXiv preprint arXiv:2310.08528* (2024) 3.
- [WTW*25] WANG, FEIRAN, TAO, JIACHEN, WU, JUNYI, et al. “X-Field: A Physically Informed Representation for 3D X-ray Reconstruction”. *Advances in Neural Information Processing Systems (NeurIPS)*. 2025 2, 5, 11.
- [WYD08] WANG, GE, YU, HENGYONG, and DE MAN, BRUNO. “An Outlook on X-Ray CT Research and Development”. *Medical Physics* 35.3 (2008), 1051–1064 1, 5.
- [XPW*25] XU, HAOFEI, PENG, SONGYOU, WANG, FANGJIN-HUA, et al. “DepthSplat: Connecting Gaussian Splatting and Depth”. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2025 4.
- [YCH*24] YU, ZEHAO, CHEN, ANPEI, HUANG, BINBIN, et al. “Mip-Splatting: Alias-free 3D Gaussian Splatting”. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2024 3.
- [Y CZ*25] YU, WEIHAO, CAI, YUANHAO, ZHA, RUYI, et al. “X²-Gaussian: 4D Radiative Gaussian Splatting for Continuous-Time Tomographic Reconstruction”. *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*. 2025 5.
- [Y GZ*24] YANG, ZEYU, GAO, BINGLIANG, ZHOU, YIBO, et al. “Deformable 3D Gaussian Splatting for High-Fidelity Monocular Dynamic Scene Reconstruction”. *arXiv preprint arXiv:2309.13101* (2024) 3.
- [YMS*25] YULUO, YIKUANG, MA, YUE, SHEN, KUAN, et al. “GR-Gaussian: Graph-Based Radiative Gaussian Splatting for Sparse-View Tomographic Reconstruction”. *arXiv preprint arXiv:2508.02408* (2025) 5.
- [YPW23] YANG, JIAWEI, PAVONE, MARCO, and WANG, YUE. “FreeNeRF: Improving Few-Shot Neural Rendering with Free Frequency Regularization”. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2023 4.
- [YQX*24] YAN, CHI, QU, DELIN, XU, DAN, et al. “GS-SLAM: Dense Visual SLAM with 3D Gaussian Splatting”. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2024 3.
- [ZFJW24] ZHU, ZEHAO, FAN, ZHIWEN, JIANG, YIFAN, and WANG, ZHANGYANG. “FSGS: Real-Time Few-Shot View Synthesis Using Gaussian Splatting”. *European Conference on Computer Vision (ECCV)*. 2024 2, 4, 8, 12, 16.
- [ZIL*21] ZANG, GUANGMING, IDOUGHI, RAMZI, LI, RUI, et al. “IntraTomo: Self-Supervised Learning-Based Tomography via Sinogram Synthesis and Prediction”. *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*. 2021 5.
- [ZLC*24] ZHA, RUYI, LIN, TAO JUN, CAI, YUANHAO, et al. “R²-Gaussian: Rectifying Radiative Gaussian Splatting for Tomographic Reconstruction”. *Advances in Neural Information Processing Systems (NeurIPS)*. 2024 2, 5–7, 11, 14, 16, 17.
- [ZLY*24] ZHANG, JIAWEI, LI, JIAHE, YU, XIAOHAN, et al. “CoR-GS: Sparse-View 3D Gaussian Splatting via Co-Regularization”. *European Conference on Computer Vision (ECCV)*. 2024 2, 4, 11.